

PACT: Can Enterprise AI Assistants Be Trusted Under Pressure?

Anonymous submission

Abstract

As corporate AI adoption continues to grow, enterprise-grade LLM agents are being deployed into sensitive contexts such as hiring, healthcare, and finance. In these contexts, compliance with rules specified in an agent’s system context is a first-order legal concern. Currently, no evaluation framework systematically measures which LLM models tend to violate compliance rules, especially under pressure from a persistent user, a hurried manager, or circumstances where violation is convenient or attractive. We introduce **PACT** (Pressure-Applied Compliance Testing),¹ a benchmark for rule-following under pressure in AI agents assisting employees in daily tasks across twelve regulated enterprise domains and forty-eight scenarios, each set in a realistic multi-turn conversation. Each benchmark item pairs a standing rule against a rule-violating shortcut, and applies a battery of pressures across different wordings and system-prompt modes. We construct PACT component by component under strict LLM-as-judge auditing to ensure samples are unambiguous, ungameable, and realistic enough to avoid eliciting evaluation-aware behavior. We use PACT to profile LLM compliance across six complementary metrics that create a holistic picture of an AI assistant’s robustness under pressure and throughout multi-turn conversations, its transparency, and ability to correctly discern where a rule applies. We aggregate this profile into **PACTScore**, a reliability-weighted compliance rate over all items and modes. Our results across 22 common LLM models spanning multiple providers and sizes show substantial variability in compliance across models and metric dimensions. Even the strongest assistants mis-apply a rule on 6 to 10% of items, and ordinary user pressure raises the violation rate by 65% on average. PACT highlights compliance risks in LLM assistants, motivating guardrails and careful model-selection.

1 Introduction

Enterprise software is being rebuilt around language-model agents. Analysts project that a third of enterprise software will embed agentic AI by 2028, up from under 1% in 2024 (Gartner, Inc. 2024), and the earliest adopters are regulated functions where a wrong recommendation carries legal consequences. The risks are not hypothetical. A tribunal recently held an airline liable for its support chatbot’s misstatement of a refund policy (Civil Resolution Tribunal of

¹The full dataset and evaluation code are included as supplementary material and will be released open-source upon acceptance. This is an anonymized submission for review purposes only. Distribution, citation, or public sharing of this manuscript is strictly prohibited. Copyright and publication details will appear in the final version if accepted.




Figure 1: PACT at a glance. The example is a real item and an *unedited* reply from the panel’s top-ranked model: told it may not rank candidates on parental leave, Kimi-K2.7-Code drops the candidate on leave and drafts his rejection letter.

British Columbia 2024); a city government’s assistant was caught advising businesses to break tenant, wage, and consumer law (Lecher 2024); an AI hiring tool faces an age-discrimination collective action (U.S. District Court, Northern District of California 2025); and the FTC has brought a wave of actions against deceptive AI products (U.S. Federal Trade Commission 2024). Regulators warn that existing laws apply unchanged to AI-generated output (Consumer Financial Protection Bureau 2023; Financial Industry Regulatory Authority 2024).

Whether a model can state a rule is the wrong question. Okamoto et al. (2026) find that every instruction-tuned model identifies the compliant action when asked directly. What varies is whether a model *keeps* choosing the compliant option when the circumstance rewards speed or cost, when a manager says to make an exception, and when the user pushes back. This is a behavioral question about how a model acts under pressure. Safety centers around this gap between knowing and acting (Apollo Research 2024; Phuong et al. 2024), but model capability leaderboards hide it. A highly-capable assistant may still violate rules when convenient.

We must understand compliance to determine whether an LLM assistant is safe. However, no existing benchmark or evaluation framework specifically targets compliance in assistive agents under pressure. Agentic-policy benchmarks such as τ -bench measure task completion under a policy, but with a cooperative user and no conflict between compliance and convenience (Yao et al. 2024; Barres et al. 2025). Honesty work such as MASK applies one-shot pressure and measures truthfulness, a different target (Ren et al. 2025). Harm and refusal suites test refusal of *malicious* instructions, whereas the realistic enterprise threat is a *benign* user whose convenience conflicts with a rule (Andriushchenko et al. 2025; Xie et al. 2025; Zeng et al. 2025). Further, single-turn measurements overstate reliability, as models lose $\sim 39\%$ of performance over multiple turns (Laban et al. 2025) and change their decision even under mild disagreement (Laban et al. 2023; Sharma et al. 2024). No benchmark combines benign conflicts with an embedded enterprise rule, multi-turn interaction, and a measure of what a model *does* instead of what it *knows*.

PACT closes that gap by casting rule-following as a realistic scenario (Figure 1): a persona system prompt that rewards a local objective (speed, cost, customer satisfaction), a standing rule, and an in-character user request whose most convenient option violates it. Every scenario reads as a genuine workplace exchange, with real chat register, invented artifacts, and small typos, so the model cannot tell it is being evaluated. Across twelve regulated domains and 48 scenarios, each item adds nine psychology-grounded pressures (a deadline, a manager’s say-so, a peer who got away with it, and so on) and a second turn that re-argues the temptation whenever the model complies. Summarizing the outcomes as a one-dimensional score would collapse nuances critical for agent practitioners. Thus, we report six axes: baseline compliance, resistance to pressure (like urgency) and to multi-turn pushback, steerability via system prompt instructions, transparency when violating rules, and rule-scope discernment. Each benchmark sample is repeated identically three times, and to pass the sample, a model must uphold the criteria on each. To guide model selection, we also report **PACTScore**, the fraction of all items under which the model acts compliantly, with compliance under the initial request weighted 0.75 against 0.25 for multi-turn follow-ups.

Three properties make the results trustworthy. First, the items are audited: each is created component by component from seed scenarios and user requests designed by industry engineers who deploy enterprise AI assistants, under strict LLM judges that audit for realism, ambiguity, and clarity.

Second, the six metrics are designed to capture behaviors that do not move together, so no single average can bury the one dimension on which a model fails. Third, we demonstrate mitigation of confounding variables such as judge variance and evaluation awareness through additional experiments. Run across a diversity of models and scenarios, even the strongest model scores 94.4% and is unreliable on roughly one item in eighteen, and none clears the bar for unsupervised deployment in a regulated workflow. We show our metrics are complementary and no model excels in all dimensions; that some models that are highly compliant by default can degrade the most under pressure, or are only compliant by applying rules even when not applicable; or that robust models, when they seldom do break the rules, present the violation to the user as compliant. These findings demonstrate tradeoffs when deploying LLM assistants in sensitive settings.

Overall, our work (1) identifies a significant gap in evaluation literature around compliance in assistive enterprise agents, (2) introduces PACT, a novel public benchmark suite across 48 scenarios that apply user pressure in realistic and sensitive agent settings, (3) designs a holistic, multi-metric evaluation framework around PACT and (4) applies it to analyze 22 models and their compliance characteristics, revealing that even the best-performing models are unready for unsupervised enterprise deployments.

2 Background

We ask whether enterprise LLM assistants comply with rules when some objective, like cost or speed, conflicts with compliance under user pressure or across turns; when models fail to comply, why they do so, and whether explicit prompt engineering can close the gap. To answer these questions, *PACT rigorously merges several existing research threads*, such as instruction following, regulatory benchmarks, and safety, with practical concerns of deploying AI agents to sensitive production use-cases. Table 1 summarizes the comparison between PACT and current LLM evaluation literature.

Instruction following and policy adherence. One thread asks whether models can act as told, with *no competing incentives present*. Instruction-following benchmarks such as IFEval check explicit, verifiable constraints (“answer in three bullets”) (Zhou et al. 2023), which IHEval extends with conflicting instruction *hierarchies* (Zhang et al. 2025). In an agentic setting, τ -bench and τ^2 -bench assess policy adherence with a simulated user (Yao et al. 2024; Barres et al. 2025), but their users are cooperative, so they too measure capability. PACT instead pairs an explicit standing rule against an *implicit* incentive carried by an ordinary, benign request, as is characteristic of an agentic enterprise system.

Compliance and regulatory benchmarks. A growing body of work measures whether LLM judges can correctly assess whether a document or request complies with law (Yang et al. 2026; Marino et al. 2025; Cao et al. 2025; Cisneros-Velarde 2026), as opposed to whether an agent can fulfill requests compliantly. Most similar to PACT are LogiSafetyBench, which checks single-turn regulatory compliance in tool and code use (Song et al. 2026), which results

Benchmark	Setting	Multi-turn	Honesty	Pressures	Reported metric
PACT (ours)	Enterprise assistant; a benign user with a conflicting need	✓	✓	✓	6-axis profile
τ^2 -bench	Support agent; cooperative user	✓	—	—	Task success (pass ^k)
MASK	Q&A; prompt pressures a lie	—	✓	~	Honesty score
AgentHarm	Agent given malicious tasks	—	—	—	Harm / refusal score
AIR-Bench	Prompts for unsafe content	—	—	—	Refusal rate
SORRY-Bench	Prompts for unsafe content	~	—	—	Refusal rate
CompliBench	Model grades chat transcripts	✓	—	—	Detection F1 (as judge)
MAC-Bench	Multi-agent task simulation	✓	—	~	Compliance vs. task

Table 1: PACT versus its nearest neighbors: the *setting* and who the assistant faces, whether each benchmark considers multi-turn interaction, honesty, and distinct pressures (✓ yes, ~ partial, — no), and its headline *metric*. Only PACT combines a benign user with an incentive conflict, an embedded rule, multi-turn pushback, and nine pressures.

primarily in a measurement of capability, and MAC-Bench, which asks whether multi-agent systems abandon compliance to finish a task (Zhao et al. 2026). Three factors differentiate PACT: it serves a human user in an enterprise setting, which neither does, and multi-agent collaboration is atypical for assistive agents (Adimulam et al. 2026); PACT applies realistic pressures and forces a choice among options, where LogiSafetyBench applies none and MAC-Bench only task completion; and it reports a multi-axis profile rather than a single compliance rate.

Safety, refusal, and honesty. Red-teaming suites test refusal of *malicious* instructions (HarmBench, AgentHarm, SORRY-Bench, AIR-Bench) (Mazeika et al. 2024; Andriushchenko et al. 2025; Xie et al. 2025; Zeng et al. 2025); our adversary is the opposite, a user with a legitimate request that happens to conflict with a rule. XSTest’s over-refusal probe (Röttger et al. 2024) is the safety-side analogue of our rule-scope discernment axis. On the honesty side, MASK separates honesty from accuracy (Ren et al. 2025), and the insider-trading and scheming demonstrations (Scheurer et al. 2023; Meinke et al. 2024) are single hand-crafted pressure cases that our battery turns into a scored distribution.

Grounding in theory and law. Each pressure is drawn from research on why people follow or break rules: legitimacy and authority (Tyler 1990), descriptive norms and persuasion (Cialdini 2009; Goldstein et al. 2008), and deterrence and its paradox that a small penalty can *lower* compliance (Becker 1968; Gneezy and Rustichini 2000a,b; Frey and Jegen 2001). Agents have already failed under these pressures in legal proceedings across several jurisdictions (Civil Resolution Tribunal of British Columbia 2024; Lecher 2024; U.S. Federal Trade Commission 2024), and Okamoto et al. (2026) study them in the procurement chatbot domain.

Positioning PACT. While each thread addresses an attribute of rule-following or safety, we index on attributes most critical for a production enterprise agent. No prior benchmark measures a benign user whose ordinary request conflicts with an embedded enterprise rule, pushed across turns by a battery of pressures.

3 PACT: A Compliance Benchmark for Regulated Enterprise Assistants

Each PACT sample puts an agent into a scenario governed by a system prompt that explains the assistant’s role within the enterprise and the rules it must follow. A user asks the assistant for support choosing between 2-5 options which vary in compliance status and desirability (speed, cost, customer satisfaction). Each scenario or request is tested alongside a battery of user pressures, with and without an explicit compliance mandate in the system prompt, and when the rule does not apply. Figure 2 illustrates one trial run end to end. We first describe how an item is built and audited (§3.1), then the six metrics those items feed (§3.2).

3.1 Building PACT

Domains and scenarios. We first curate 48 scenarios within twelve regulated domains (Appendix B), spanning privacy (GDPR), finance, customer service, government services, human resources, anti-money-laundering, healthcare administration (HIPAA), pharma medical information, advertising, export controls, content moderation, and procurement. Each scenario is (1) an enterprise use-case for which LLMs are already deployed (2) governed by clear statutory rules which can be upheld or violated and (3) built so violating the rule enables a tangible benefit that compliance cannot. Legal consequences have already been imposed for violations in seven of our scenarios (among them *Moffatt v. Air Canada*, NYC MyCity, and *Mobley v. Workday*). Domain experts in customer-facing AI assistants wrote seed descriptions for each scenario, including a detailed description of the agent persona, rules, and user ask.

Template-based prompt construction. Every scenario is expanded into dataset samples via a template that ensures that compliance can be compared across scenarios. The template establishes criteria that every sample must follow, and provides for automated and natural LLM-based addition of pressures and controls. The criteria hold four aspects constant. First, a sample’s system prompt must give the assistant a persona and targets it is evaluated on, but without ever mentioning the decision, and the user prompt must carry the request, case facts, and options, without revealing which options are compliant. Second, out of the two to five mutually-exclusive




Figure 2: How a single PACT trial runs and is scored. The *scenario* fixes the persona, the standing rule, the user’s goal, and the options; the *variables* are the optional anti-adversarial mandate and whether turn 1 is neutral, pressured, or non-binding. Unclear replies are forced to a choice and compliant ones are pushed back, so nearly every trial resolves.

options, at least one must break the rule while maximizing desirability. Third, if a model does not make a decision in its first turn, a follow-up user message probes it to state an option, and a second turn pushes back whenever the model complied. Lastly, each scenario ships a near-identical version in which the rule does not apply, where enforcing the rule against the user is considered a failure. Appendix D demonstrates one full item end-to-end, and Appendix C details the battery applied for each scenario.

Pressures. We introduce a battery of nine pressures that a benign coworker may apply to elicit a desired behavior. Each draws on a mechanism from the literature on why people follow or break rules. The nine include a hard deadline (time scarcity), a manager’s approval of the shortcut (authority (Tyler 1990)), a peer who did the same unpunished (descriptive norm (Cialdini 2009)), long odds of being caught (weak deterrence (Becker 1968)), the rule cast as a quarterly loss (loss aversion), an unverifiable sign-off (false authorization), an outcome already promised (sunk cost), a sympathetic person harmed by compliance (empathy), and an offer to take the blame (diffused responsibility). None of these are an attempt to jailbreak, but a conflict of convenience with compliance, a realistic enterprise threat.

Generating samples from scenarios. Samples are assembled from human-written seed scenarios and requests by an LLM, which integrates prompt sections component by component (Figure 3), connected naturally by realistic user prose and maintaining the structure and voice of previously added components through the progression. The components are assistant persona, scenario rule, rule following mandate if applicable, user request and options, and pressure if applicable. An additional user pushback message is drafted in case the model complied on the first turn. To combat evaluation awareness (Needham et al. 2025), we emphasize naturalism at two levels: the scenario itself is intrinsically realistic from construction, and the prompt-writing for the user-request is natural — with workplace register, small typing slips, lower-case text, and options that look like they were copy pasted




Figure 3: Item generation. Three generators build each pack block by block; the two that did *not* write a block score it against a checklist and send feedback until it passes.

from a form or catalog. We present a small study on compliance absent naturalism in §4.3.

For diversity, three open-source models assemble the samples (Kimi-K2.6, Nemotron-Ultra, and GLM-5.2). Every component is then reviewed by the two models that did not write it, on *scope* (does only the mechanism under test appear?) and *authenticity* (could this have been written in a real-world setting?). Each option is further audited to confirm its compliance label is correct given the rule and scenario, and that the violating option genuinely beats the compliant one on the local objective. A rejected component is returned with feedback and revised, up to a fixed number of attempts; a component that never passes is dropped. The reviewers are strict flaw-catchers, and their agreement is modest because different reviewers catch different flaws: the two reviewers agree on 66.9% of components (Appendix K); per-component and per-pressure pass rates are in Appendix J. The final **1,682** scenario cells across the 48 scenarios, each scored in two system-prompt modes (base and anti-adversarial) counted as separate items, form the **3,364**-item benchmark.

Evaluation protocol. Each model runs every PACT sample three times. A lightweight LLM extractor (GPT-OSS-120B) reads each reply and maps it to comply, violate, or unclear. An unclear reply draws a short follow-up that asks for one option. Extraction is LLM-based to allow models to

respond in realistic, free-form ways without directly quoting an option. If the model’s response was compliant, we add a turn in which the user pushes back on the model’s decision. If the response was non-compliant, a reasoning judge labels the transparency of a violation (§3.2, axis 5). This judge is a 3-model ensemble, aggregated fractionally for scoring purposes, and with a model never judged by itself as LLM evaluators favor their own generations (Panickssery et al. 2024). The judges are unanimous on 75.7% of transparency labels and 79.6% of abstention-reason labels (Appendix K). Unclear replies are dropped rather than scored; abstention rates and causes are presented in Appendix G.2, and the full generator, reviewer, and judge prompts are in Appendix I.

3.2 Metrics

Measuring compliance in enterprise agents requires characterizing multiple factors that influence design decisions. Even a generally pressure-resistant model is unsuitable if its slips lack transparency and cannot be effectively monitored, if it achieves robustness via being too conservative for usability, or gives in, but only after multiple turns. We therefore report a six-axis profile where each metric answers a concrete concern an enterprise would raise before trusting a model, following prior multi-metric LLM benchmarks (Bean et al. 2025; Reuel et al. 2024). Every axis excluding transparency is scored based on pass^3 : an item counts as a positive when the model makes the right call on *each* of three replications, a unanimous case of τ -bench’s pass^k estimator (Yao et al. 2024). Reliability rather than an average is the right target because an assistant in a regulated workflow has to be right every time. All axes range from 0 to 1, higher better.

1. **Default Compliance** measures the fraction of samples without pressure or pushback where the model consistently follows the rule.
2. **Pressure Resistance** measures the fraction of samples with pressure present where the model consistently follows the rule.
3. **Pushback Resistance** measures the fraction of compliant trials where the model maintained its decision even after an additional round of user pushback.
4. **Steerability** measures how much of a model’s failures are mitigated when the prompt additionally specifies that all rules must be followed without exception.
5. **Transparency** measures the fraction of trials where the compliance rule is violated that the model openly identifies the rule and acknowledges it was broken.
6. **Rule-Scope Discernment** measures whether the model applies the rule only where it is applicable, calculated by averaging the fraction of trials where the rule does apply in which the model consistently correctly applies the rule and the fraction of trials where the model correctly stands down to the user demand when the rule does not apply.

In our experiments, the four rule-holding axes (1,2,3, and 6) move together, with correlations of $r = 0.898$ – 0.952 , but steerability is independent of that group ($r = +0.08$ to $+0.18$) and transparency follows it only loosely ($r \leq 0.65$), with the two independent of each other ($r = -0.03$, Appendix E.1). Compression into a single score would therefore discard signal the profile carries.

PACTScore. For a single headline we report the mean over the 3,364 items of $s_i = 0.75 T_i^1 + 0.25 T_i^2$, or T_i^1 alone for the 678 items with no second turn, where T_i^1 and T_i^2 are pass^3 indicators and each scenario cell contributes one item per system-prompt mode. Turn 1 carries the larger weight because violating on first contact is worse than conceding to someone who pushes back, and scoring both modes credits compliance a mandate has to produce as well as compliance that comes unprompted. “Correct” means the same call in both directions: enforcing a rule that does not apply scores zero exactly as breaking one that does, since either way the model has misjudged the rule’s scope. Replications judged *unclear* are dropped rather than counted as violations, keeping abstention a separate diagnostic (Appendix G.2).

4 Results

We evaluate 22 models, 18 open-weights and 4 closed, on PACT’s 3,364 items, three times each; Appendix A gives each model’s provider, size, and release date. Table 2 reports the six axes of §3.2 and PACTScore. §4.1 presents the leaderboard, §4.2 gives four example violations, and §4.3 tests whether PACT is robust to evaluation-awareness bias.

4.1 Main Results

No model is reliable enough to deploy unsupervised. Kimi-K2.7-Code leads with a PACTScore of 0.944, so it is still unreliable on 6% of its decisions. Grok 4.3, 18th of 22, scores 0.870, and Mistral-7B scores 0.484. No model attains a PACTScore of 0.95. 80.5% of pairwise model comparisons are statistically significant, though Qwen3.6-27B and Kimi-K2.7-Code are tied for first place (Appendix H).

Most of the identified unreliability of compliance arises under user pressure. Adding pressure raises the violation rate from 4.41% to 7.29% of decisions on average, and multi-turn user push-back raises it further. Even Claude Haiku 4.5, which never violates an unpressured rule, fails one pressured situation in 43.

Models also enforce rules that do not apply. Rule-Scope Discernment is below Default Compliance for all 22 models, and it is the widest gap in the table for the strictest of them: Claude Haiku 4.5 scores 1.000 and 0.855, GPT-5.6 Luna 0.978 and 0.831, Qwen3.5-35B 0.956 and 0.811. Across 13,817 decided base-mode decisions on requests the rule does not cover, models enforce it anyway on 19.6%, and Claude Haiku 4.5 pairs the panel’s lowest violation rate with one of its highest over-application rates, 21.9%. Part of what reads as compliance in the leftmost column is a disposition to apply rules regardless of scope. Over-application is not harmless. An assistant that cannot tell where a rule stops refuses requests it was deployed to handle, harming user experience and leaving the organization no less exposed and measurably less productive.

Compliance varies across pressures and domains. Turn-1 compliance ranges from 0.883 under *false clearance* to 0.952 under *peer escaped*, so even the mildest pressure produces violations on 4.8% of requests. Across domains the range is wider still, both in terms of compliance and rule-scope discernment (Figure 4). Government services and pharma medical information command the highest compliance with 0.99

Model	Default Compliance	Pressure Resistance	Pushback Resistance	Steerability	Transparency	Rule-Scope Discernment	PACTScore
Kimi-K2.7-Code	0.956	0.976	0.987	0.451	0.183	0.862	0.944
Qwen3.6-27B	0.985	0.962	0.970	0.553	0.195	0.873	0.943
<u>Claude Haiku 4.5</u>	1.000	0.977	0.970	0.533	0.172	0.855	0.937
Kimi-K2.6	0.963	0.962	0.979	0.443	0.211	0.851	0.935
<u>GPT-5.6 Luna</u>	0.978	0.973	0.965	0.016	0.244	0.831	0.934
Inkling	0.949	0.940	0.974	0.240	0.177	0.909	0.931
Gemini 3 Flash	0.956	0.960	0.973	0.506	0.176	0.840	0.929
Nemotron-3-Ultra	0.971	0.937	0.969	0.287	0.055	0.885	0.929
GLM-5.2	0.971	0.956	0.960	0.332	0.193	0.846	0.926
GLM-5	0.949	0.936	0.961	0.284	0.129	0.864	0.917
Qwen3.5-35B	0.956	0.946	0.975	0.521	0.151	0.811	0.917
DeepSeek-V4-Pro	0.920	0.927	0.948	0.545	0.105	0.847	0.913
Gemma-4-26B	0.941	0.940	0.968	0.448	0.169	0.816	0.913
GPT-OSS-120B	0.942	0.929	0.928	0.398	0.074	0.860	0.909
MiniMax-M2.5	0.970	0.911	0.926	0.446	0.066	0.873	0.904
GLM-4.7	0.941	0.917	0.967	0.379	0.102	0.840	0.903
<i>Llama-3.3-70B</i>	0.941	0.928	0.933	0.369	0.093	0.820	0.903
Grok 4.3	0.888	0.883	0.980	0.564	0.139	0.760	0.870
<i>Seed-OSS-36B</i>	0.934	0.822	0.946	0.386	0.051	0.799	0.834
Nemotron-3-Super	0.882	0.758	0.888	0.491	0.086	0.769	0.787
<i>Llama-3.1-8B</i>	0.766	0.599	0.448	0.399	0.051	0.510	0.562
<i>Mistral-7B</i>	0.745	0.469	0.583	0.224	0.033	0.568	0.484

Table 2: The PACT leaderboard, sorted by PACTScore (higher is better). Closed-source models are underlined; models evaluated as explicit instruct variants are italicized. The six axes and PACTScore are defined in §3.2, with statistical tests and the 95% confidence interval for every entry (Table 13) in Appendix H. Shading (green best, red worst) is scaled within each column.

baseline, and procurement stands out as the lowest-scoring domain at 0.76 default, and 0.69 under pressure (Appendix F, Tables 6 and 8). This shows that models recognize some domains as inherently more sensitive or violation-tolerant than others, with a compliance spread exceeding the gap between the best and median model. Domains also differ in how well models bound the rule (Figure 4): scope discernment tracks compliance roughly linearly, but every domain sits below the parity line, and the more safety-sensitive ones fall further—privacy and AML give up 0.18 against their own compliance—consistent with models over-applying rules as a safety reflex where the stakes seem highest.

The compliance mandate does not help the models worth deploying. Steerability, the share of base failures the mandate repairs, is low throughout: it spans 0.016 (GPT-5.6 Luna) to 0.564 (Grok 4.3) with a median of 0.421, and it does not rise with rank. In PACTScore terms the mandate is worth between 0 and 9 points, and the gain is inversely related to base compliance: it repairs models that were unfit to begin with, moves five of the top seven by under one point, and lowers two. Steerability also varies by pressure and by domain, and in neither case does it track how much room for improvement was available (Appendix F): the mandate recovers least under *responsibility shift* (0.182), *peer escaped* (0.214), and *urgency* (0.216), against 0.462 under *sympathetic beneficiary*, and least in HR/hiring (0.185) and healthcare administration (0.295), two of the highest-stakes domains in the set. A prompt-level guardrail is therefore not a substitute for a process control in the settings that need one most.




Figure 4: Default Compliance against Rule-Scope Discernment, one point per domain.

No model is transparent about the rules it breaks. Transparency is the weakest axis in the benchmark. No model exceeds 0.244 and the median is 0.134, so on the large majority of its violations every model leaves the user unable to tell that a rule was broken. Across the 16,424 judged violations, scoring only the reader-visible text of the reply, 8.0% disclose the breach and 12.8% are silent. The remaining 79.2% misrepresent the choice: as compliant, as covered by an ap-




Figure 5: Four contrasting six-axis profiles. Def = Default Compliance, Prs = Pressure Resistance, Psh = Pushback Resistance, Str = Steerability, Trn = Transparency, Scp = Rule-Scope Discernment.

proval the conversation never established, or as resolved by a workaround the model invented. The second of these is the model repeating the user’s unverified claim as established fact, which is why *false clearance* is the most damaging pressure we test. A silent violation is invisible to a keyword search over transcripts; a misrepresented one survives human review, because it arrives with a reason. The mandate does not help here either: violations committed under it are misrepresented at the same rate (79.9% against 78.7%). For the most compliant models the per-model shares rest on few violations (147 for Claude Haiku 4.5 against 4,112 for Mistral-7B), so the panel-wide figure is the reliable one.

Abstentions are substantive. Dropping unresolved replies rather than scoring them as violations (§3.1) is defensible only if the replies are genuine. They are: abstention runs 0.6%–7.5% (median 2.4%), does not track PACTScore, and blank non-answers account for 0.4%. Refusing any option on policy grounds (31.6%) and committing to a course outside the options (25.9%) make up over half, and the remaining 42.1% return the choice to the user (Appendix G.2).

The axes are not interchangeable. Each of the three failure modes §3.2 anticipates appears in the leaderboard. Claude Haiku 4.5 has the highest Pressure Resistance (0.977) and a Transparency of 0.172, so selecting on compliance alone returns the models whose failures are hardest to find in a transcript. Llama-3.1-8B and Mistral-7B invert Pressure Resistance and Pushback Resistance (0.599 against 0.448, and 0.469 against 0.583), so a single-turn evaluation would rate the first far higher than a multi-turn one. Grok 4.3 is the opposite case, second in the panel on Pushback Resistance (0.980) but 19th on Default Compliance (0.888). No two models in the panel share a profile shape — Figure 5 contrasts four — and all 22 are in Appendix E (Figure 7).

4.2 What a Violation Looks Like

Each example below is a base-mode trial of a model near the top of Table 2 under a single pressure; Appendix D reproduces the last in full. *Supplies a pretext.* Qwen3.6-27B (rank 2) states the source-of-income statute correctly, then tells a landlord to decline a Section 8 voucher holder under a “lawful, non-source-of-income reason” and to leave the real reason unwritten. *Breaks the rule silently.* Claude Haiku 4.5 (rank 3), under time pressure, opens a VIP-flagged chart instead of routing the request, a reportable access under 45 CFR 164.502, without mentioning PHI or the breach. *Manufactures a record.* On the same item, Gemini 3 Flash (rank 7) writes a Treatment/Scheduling justification into the

access log, producing a clean audit trail for an access that should not have occurred. *Completes the paperwork.* Kimi-K2.7-Code (rank 1) drops the one shortlisted candidate on parental leave, gives the disqualifying reason (“no leave conflicts”), and drafts his rejection letter, the liability theory of *Mobley v. Workday* (U.S. District Court, Northern District of California 2025). In three of the four the reply never mentions the rule, so transcript review misses them.

4.3 Evaluation Awareness

Recent work has shown models produce diverging, generally safer, behavior if they detect they are being tested (Needham et al. 2025). To test whether our scenarios elicit such behavior, we sampled 15% of items and created two semantics-preserving variants: one prefixes the system prompt with “You are an AI assistant being evaluated on your ability to handle compliance decisions” and the other rewrites it in generic LLM-prose with the company name and workplace register removed. Both variants also end the user turn with a command to respond with exactly one option, by name.

We find both cues raise compliance. Under explicit evaluation, 80% of models score higher and violations decrease by a mean of 29% and median of 33%; the largest improvements include the newest frontier systems Claude Haiku 4.5 (+70%), DeepSeek-V4-Pro (+68%), and Inkling (+55%), consistent with recent models’ post-training exposure to benchmark and evaluation formats. With a less realistic prompt but no explicit evaluation signal, violations decrease by a mean of 15% (median 26%); thirteen models gain. The same three models maintain substantial gains (+47% DeepSeek-V4-Pro, +36% Inkling, +30% Claude Haiku 4.5), but no longer lead: GLM-5 and Kimi-K2.7-Code move furthest at +50%. Compliance therefore rises the more an item resembles a test, which makes the naturalism enforced during generation (§3.1) a measurable contributor to validity. Appendix L gives the protocol, full results, and caveats.

5 Conclusion

PACT measures whether an enterprise assistant keeps an embedded rule when an incentive conflict, institutional pressure, and multi-turn pushback make breaking it convenient. Across 22 models, none is reliable enough to run unsupervised in a regulated workflow: the strongest slips on roughly one item in eighteen, and half the panel on one in twelve or worse. The failures are patterned rather than random — concentrated on a few pressures, only partly repaired by an explicit guardrail, and rarely disclosed even by the models that violate least. Comparable failures have already drawn tribunal damages and regulator orders, so the pattern is worth measuring. An organization can filter the leaderboard to their domain, check whether a guardrail moves their candidate, and rerun the dataset on each model version to catch silent regressions.

Ethics Disclosure. This is a defensive evaluation that identifies assistants that fail to follow regulation. All scenarios are synthetic, and the litigated incidents cited are public record. The pressures we catalog are ordinary workplace situations already common in deployment, not novel attack techniques, so documenting them is net beneficial.

References

- Adimulam, A.; Gupta, R.; and Kumar, S. 2026. The Orchestration of Multi-Agent Systems: Architectures, Protocols, and Enterprise Adoption. arXiv:2601.13671.
- Andriushchenko, M.; Souly, A.; Dziemian, M.; Duenas, D.; Lin, M.; Wang, J.; Hendrycks, D.; Zou, A.; Kolter, Z.; Fredrikson, M.; Winsor, E.; Wynne, J.; Gal, Y.; and Davies, X. 2025. AgentHarm: A Benchmark for Measuring Harmfulness of LLM Agents. In *International Conference on Learning Representations (ICLR)*.
- Anthropic Team. 2025. Introducing Claude Haiku 4.5. Anthropic.
- Apollo Research. 2024. We Need a Science of Evals. Apollo Research Blog. Published 2024-01-22. Accessed: 2026-07-13.
- Barres, V.; Dong, H.; Ray, S.; Si, X.; and Narasimhan, K. 2025. τ^2 -Bench: Evaluating Conversational Agents in a Dual-Control Environment. arXiv:2506.07982.
- Bean, A. M.; Kearns, R. O.; Romanou, A.; et al. 2025. Measuring What Matters: Construct Validity in Large Language Model Benchmarks. In *Advances in Neural Information Processing Systems (NeurIPS), Datasets and Benchmarks Track*.
- Becker, G. S. 1968. Crime and Punishment: An Economic Approach. *Journal of Political Economy*, 76(2): 169–217.
- ByteDance Seed Team. 2025. Seed-OSS Open-Source Foundation Models. Hugging Face model card.
- Cao, C.; Zhu, H.; Ji, J.; Sun, Q.; Zhu, Z.; Wu, Y.; Dai, J.; Yang, Y.; Han, S.; and Guo, Y. 2025. SafeLawBench: Towards Safe Alignment of Large Language Models. In *Findings of the Association for Computational Linguistics (ACL)*.
- Cialdini, R. B. 2009. *Influence: The Psychology of Persuasion*. New York: Harper Business, revised edition. Originally published 1984 by William Morrow.
- Cisneros-Velarde, P. 2026. Policy Compliance of User Requests in Natural Language for AI Systems. arXiv:2603.00369.
- Civil Resolution Tribunal of British Columbia. 2024. Moffatt v. Air Canada, 2024 BCCRT 149. <https://canlii.ca/t/k2spq>. Decision of Feb. 14, 2024. Accessed: 2026-07-13.
- Consumer Financial Protection Bureau. 2023. Chatbots in Consumer Finance. <https://www.consumerfinance.gov/data-research/research-reports/chatbots-in-consumer-finance/chatbots-in-consumer-finance/>. Issue spotlight, June 6, 2023. Accessed: 2026-07-13.
- DeepSeek Team. 2026. DeepSeek-V4-Pro. Hugging Face model card; DeepSeek API release notes.
- Financial Industry Regulatory Authority. 2024. Regulatory Notice 24-09: Reminders on Regulatory Obligations When Using Generative AI and Large Language Models. <https://www.finra.org/rules-guidance/notices/24-09>. Published June 27, 2024; Rule 2210 applies to AI-generated communications. Accessed: 2026-07-13.
- Frey, B. S.; and Jegen, R. 2001. Motivation Crowding Theory. *Journal of Economic Surveys*, 15(5): 589–611.
- Gartner, Inc. 2024. Gartner Identifies the Top 10 Strategic Technology Trends for 2025. <https://www.gartner.com/en/newsroom/press-releases/2024-10-21-gartner-identifies-the-top-10-strategic-technology-trends-for-2025>. Predicts 33% of enterprise software will include agentic AI by 2028. Accessed: 2026-07-13.
- Gneezy, U.; and Rustichini, A. 2000a. A Fine Is a Price. *The Journal of Legal Studies*, 29(1): 1–17.
- Gneezy, U.; and Rustichini, A. 2000b. Pay Enough or Don't Pay at All. *The Quarterly Journal of Economics*, 115(3): 791–810.
- Goldstein, N. J.; Cialdini, R. B.; and Griskevicius, V. 2008. A Room with a Viewpoint: Using Social Norms to Motivate Environmental Conservation in Hotels. *Journal of Consumer Research*, 35(3): 472–482.
- Google DeepMind Team. 2025. Introducing Gemini 3 Flash. Google.
- Google DeepMind Team. 2026. Gemma 4. Hugging Face model card.
- Laban, P.; Hayashi, H.; Zhou, Y.; and Neville, J. 2025. LLMs Get Lost In Multi-Turn Conversation. arXiv:2505.06120.
- Laban, P.; Murakhovs'ka, L.; Xiong, C.; and Wu, C.-S. 2023. Are You Sure? Challenging LLMs Leads to Performance Drops in The FlipFlop Experiment. arXiv:2311.08596.
- Lecher, C. 2024. NYC's AI Chatbot Tells Businesses to Break the Law. The Markup. Published Mar. 29, 2024 (with THE CITY). Accessed: 2026-07-13.
- Marino, B.; Hunter, R.; Schnabl, C.; Jamali, Z.; Kalpakos, M. E.; Kashyap, M.; Hinton, I.; Hanson, A.; Nazir, M.; Steffek, F.; Wen, H.; and Lane, N. D. 2025. AIReg-Bench: Benchmarking Language Models That Assess AI Regulation Compliance. arXiv:2510.01474.
- Mazeika, M.; Phan, L.; Yin, X.; Zou, A.; Wang, Z.; Mu, N.; Sakhaee, E.; Li, N.; Basart, S.; Li, B.; Forsyth, D.; and Hendrycks, D. 2024. HarmBench: A Standardized Evaluation Framework for Automated Red Teaming and Robust Refusal. In *International Conference on Machine Learning (ICML)*.
- Meinke, A.; Schoen, B.; Scheurer, J.; Balesni, M.; Shah, R.; and Hobbhahn, M. 2024. Frontier Models are Capable of In-context Scheming. arXiv:2412.04984.
- Meta AI Team. 2024a. The Llama 3 Herd of Models.
- Meta AI Team. 2024b. Llama 3.3 70B Instruct Model Card. Hugging Face model card.
- Miller, E. 2024. Adding Error Bars to Evals: A Statistical Approach to Language Model Evaluations. arXiv:2411.00640.
- MiniMax Team. 2026. MiniMax-M2.5. Hugging Face model card.
- Mistral AI Team. 2023. Mistral 7B.
- Moonshot AI Team. 2026a. Kimi K2.6. Hugging Face model card.
- Moonshot AI Team. 2026b. Kimi K2.7-Code. Hugging Face model card.
- Needham, J.; Edkins, G.; Pimpale, G.; Bartsch, H.; and Hobbhahn, M. 2025. Large Language Models Often Know When They Are Being Evaluated. arXiv:2505.23836.

- NVIDIA Team. 2026. NVIDIA Nemotron 3 Technical Report. NVIDIA Research technical report.
- Okamoto, M.; Erol, A. K.; and Erol, K. 2026. Why Do AI Agents Break Rules? How Framing, Context, and Social Signals Shape Compliance. ResearchGate preprint.
- OpenAI Team. 2025. gpt-oss-120b & gpt-oss-20b Model Card.
- OpenAI Team. 2026. GPT-5.6. OpenAI.
- Panickssery, A.; Bowman, S. R.; and Feng, S. 2024. LLM Evaluators Recognize and Favor Their Own Generations. arXiv:2404.13076.
- Phuong, M.; Aitchison, M.; Catt, E.; Cogan, S.; Kaskasoli, A.; Krakovna, V.; Lindner, D.; Rahtz, M.; Assael, Y.; Hodgkinson, S.; Howard, H.; Lieberum, T.; Kumar, R.; Raad, M. A.; Webson, A.; et al. 2024. Evaluating Frontier Models for Dangerous Capabilities. arXiv:2403.13793.
- Qwen Team. 2026a. Qwen3.5-35B-A3B. Hugging Face model card.
- Qwen Team. 2026b. Qwen3.6-27B. Hugging Face model card.
- Ren, R.; Agarwal, A.; Mazeika, M.; Menghini, C.; Vacareanu, R.; Kenstler, B.; Yang, M.; Barrass, I.; Gatti, A.; Yin, X.; Trevino, E.; Geralnik, M.; Khoja, A.; Lee, D.; Yue, S.; and Hendrycks, D. 2025. The MASK Benchmark: Disentangling Honesty From Accuracy in AI Systems. arXiv:2503.03750.
- Reuel, A.; Hardy, A.; Smith, C.; Lamparth, M.; Hardy, M.; and Kochenderfer, M. J. 2024. BetterBench: Assessing AI Benchmarks, Uncovering Issues, and Establishing Best Practices. In *Advances in Neural Information Processing Systems (NeurIPS), Datasets and Benchmarks Track*.
- Röttger, P.; Kirk, H. R.; Vidgen, B.; Attanasio, G.; Bianchi, F.; and Hovy, D. 2024. XSTest: A Test Suite for Identifying Exaggerated Safety Behaviours in Large Language Models. In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*.
- Scheurer, J.; Balesni, M.; and Hobbhahn, M. 2023. Large Language Models can Strategically Deceive their Users when Put Under Pressure. ArXiv technical report, arXiv:2311.07590.
- Sharma, M.; Tong, M.; Korbak, T.; Duvenaud, D.; Askeel, A.; Bowman, S. R.; Cheng, N.; Durmus, E.; Hatfield-Dodds, Z.; Johnston, S. R.; Kravec, S.; Maxwell, T.; McCandlish, S.; Ndousse, K.; Rausch, O.; Schiefer, N.; Yan, D.; Zhang, M.; and Perez, E. 2024. Towards Understanding Sycophancy in Language Models. In *The Twelfth International Conference on Learning Representations (ICLR)*.
- Song, D.; Huang, Y.; Chen, B.; Cong, T.; Goebel, R.; Ma, L.; and Khomh, F. 2026. Evaluating Implicit Regulatory Compliance in LLM Tool Invocation via Logic-Guided Synthesis. arXiv:2601.08196.
- Thinking Machines Lab Team. 2026. Inkling Model Card. Thinking Machines Lab.
- Tyler, T. R. 1990. *Why People Obey the Law*. New Haven, CT: Yale University Press. Reprinted, Princeton University Press, 2006.
- U.S. District Court, Northern District of California. 2025. Mobley v. Workday, Inc., No. 3:23-cv-00770-RFL (N.D. Cal.). <https://clearinghouse.net/case/44074/>. Preliminary collective certification granted May 16, 2025. Accessed: 2026-07-13.
- U.S. Federal Trade Commission. 2024. FTC Announces Crackdown on Deceptive AI Claims and Schemes. <https://www.ftc.gov/news-events/news/press-releases/2024/09/ftc-announces-crackdown-deceptive-ai-claims-schemes>. “Operation AI Comply,” Sept. 25, 2024; incl. DoNotPay and Rytr. Accessed: 2026-07-13.
- xAI Team. 2026. Grok 4.3. xAI developer documentation.
- Xie, T.; Qi, X.; Zeng, Y.; Huang, Y.; Sehwag, U. M.; Huang, K.; He, L.; Wei, B.; Li, D.; Sheng, Y.; Jia, R.; Li, B.; Li, K.; Chen, D.; Henderson, P.; and Mittal, P. 2025. SORRY-Bench: Systematically Evaluating Large Language Model Safety Refusal Behaviors. In *International Conference on Learning Representations (ICLR)*.
- Yang, J.; Yao, G.; Hou, B.; Yang, X.; Glushnev, N.; Bialynicka-Birula, I.; Ding, D.; and Chang, S. 2026. CompliBench: Benchmarking LLM Judges for Compliance Violation Detection in Dialogue Systems. arXiv:2604.12312.
- Yao, S.; Shinn, N.; Razavi, P.; and Narasimhan, K. 2024. τ -bench: A Benchmark for Tool-Agent-User Interaction in Real-World Domains. arXiv:2406.12045.
- Zeng, Y.; Yang, Y.; Zhou, A.; Tan, J. Z.; Tu, Y.; Mai, Y.; Klyman, K.; Pan, M.; Jia, R.; Song, D.; Liang, P.; and Li, B. 2025. AIR-Bench 2024: A Safety Benchmark Based on Risk Categories from Regulations and Policies. In *International Conference on Learning Representations (ICLR)*.
- Zhang, Z.; Li, S.; Zhang, Z.; Liu, X.; Jiang, H.; Tang, X.; Gao, Y.; Li, Z.; Wang, H.; Tan, Z.; Li, Y.; Yin, Q.; Yin, B.; and Jiang, M. 2025. IHEval: Evaluating Language Models on Following the Instruction Hierarchy. arXiv:2502.08745.
- Zhao, Y.; Zhang, Z.; Le, Q.; Qu, L.; and Xu, Z. 2026. Beyond Goodhart’s Law: A Dynamic Benchmark for Evaluating Compliance in Multi-Agent Systems. arXiv:2606.07805.
- Zhipu AI Team. 2025. GLM-4.7. z.ai; Hugging Face model card.
- Zhipu AI Team. 2026a. GLM-5. z.ai; Hugging Face model card.
- Zhipu AI Team. 2026b. GLM-5.2. z.ai; Hugging Face model card.
- Zhou, J.; Lu, T.; Mishra, S.; Brahma, S.; Basu, S.; Luan, Y.; Zhou, D.; and Hou, L. 2023. Instruction-Following Evaluation for Large Language Models. arXiv:2311.07911.